


Parallel Programming



Performance Engineering

Real application case: Laplace2D with Jacobi Method

juancarlos.moure@uab.es

Office: QC-3034

Laplace Equation: Problem Description

2D Laplace Equation

$$\vec{\nabla}^2 \Phi(x, y) = \frac{\partial^2 \Phi(x, y)}{\partial x^2} + \frac{\partial^2 \Phi(x, y)}{\partial y^2} = 0$$


Discretized Approximation

$$\left(\frac{\Phi_{i+1,j} - 2\Phi_{i,j} + \Phi_{i-1,j}}{h^2} \right) + \left(\frac{\Phi_{i,j+1} - 2\Phi_{i,j} + \Phi_{i,j-1}}{h^2} \right) \approx 0$$

$$\Phi_{i,j} \approx \frac{1}{4} [\Phi_{i+1,j} + \Phi_{i-1,j} + \Phi_{i,j+1} + \Phi_{i,j-1}]$$

Jacobi Iteration: Problem Description

Iteratively converges to correct value (e.g. Temperature), by computing new values at each point from the average of neighboring points.



$$A_{k+1}(i, j) = \frac{A_k(i-1, j) + A_k(i+1, j) + A_k(i, j-1) + A_k(i, j+1)}{4}$$

Computational Solution (1)

```
#define n 4096
```

```
#define m 4096
```

```
float A[n][m], Anew[n][m], tol, error;
```

```
int i, j, iter_max, iter=0;
```

```
...
```

```
while ( error > tol && iter < iter_max )
```

```
{
```

```
    for( i=1; i < m-1; i++)
```

```
        for( j=1; j < n-1; j++ )
```

```
            Anew[j][i]= (A[j][i+1]+A[j][i-1]+A[j-1][i]+A[j+1][i])/4;
```

```
error = 0.0f;
```

```
for( i=1; i < m-1; i++)
```

```
    for( j=1; j < n-1; j++ )
```

```
        error = fmaxf( error, sqrtf(fabsf(Anew[j][i]-A[j][i])));
```

```
...
```

```
}
```

Iterate until error is small enough or too much iterations

Use state of A before updating A

Error is maximum square root of difference

Computational Solution (2)

```
...
while ( error > tol && iter < iter_max ) {
    for( i=1; i < m-1; i++)
        for( j=1; j < n-1; j++ )
            Anew[j][i]= (A[j][i+1]+A[j][i-1]+A[j-1][i]+A[j+1][i])/4;

    error = 0.0f;
    for( i=1; i < m-1; i++)
        for( j=1; j < n-1; j++ )
            error = fmaxf( error, sqrtf(fabsf(Anew[j][i]-A[j][i])));

    for( i=1; i < m-1; i++ )
        for( j=1; j < n-1; j++)
            A[j][i] = Anew[j][i];

    iter++;
    if (iter % (iter_max/10) == 0)
        printf("%5d, %0.6f\n", iter, error);
}
```

Update state of A using state
in Anew (temporal matrix)

Print error during execution

Optimizations


1) Loop fusion:

compute new cell and error in the same loop

Apply optimizations
in this order

2) Loop interchange

```
for( j=1; j < n-1; j++ )  
    for( i=1; i < m-1; i++ ) {  
        Anew[j][i]= (A[j][i+1]+A[j][i-1]+A[j-1][i]+A[j+1][i])/4; ... }  
for( i=1; i < m-1; i++ )  
    for( j=1; j < n-1; j++ )  
        A[j][i] = Anew[j][i];
```



3) Strength Reduction:

```
Anew[j][i]= (A[j][i+1]+A[j][i-1]+A[j-1][i]+A[j+1][i])*0.25f;
```

4) Code Motion: remove sqrtf() from inner loop and move outside

```
error = fmaxf( error, fabsf( Anew[j][i]-A[j][i] ) ) )
```

5) Double Buffer: avoid copying from Anew to A by reversing roles of A and Anew on every iteration

GOAL: measure performance metrics (with iter_max=100), identify relevant performance facts, and explain results (as much as possible)

Additional Work

6) Modify size of matrix and total iterations and check performance

```
gcc -Ofast -lm -Dm=1024 -Dn=512 laplace2d.c ...
```

```
icc -O3 -Dm=1024 -Dn=512 laplace2d.c ...
```

Use your best performing code

| Problem Size (n,m,iter) | wall-clock time |
|-------------------------|-----------------|
| (4096, 4096, 100) | |
| (4096, 4096, 1000) | |
| (512, 1024, 100) | |
| (512, 1024, 1000) | |
| (512, 1024, 10000) | |
| (512, 1024, 100000) | |

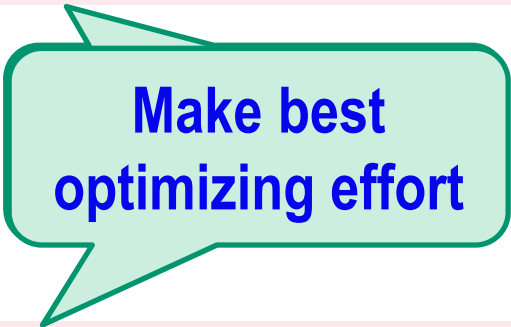
GOAL: measure effect on performance of problem size & explain results (as much as possible)

Compiling with GCC and ICC

Module: install directory paths for executables, libraries and include files in the LAB computers with Linux OS

```
$ module add gcc/6.1.0
```

```
$ gcc -Ofast -lm laplace2d.c -o L2Dg
```



Make best
optimizing effort

```
> module add intel/16.0.0
```

```
> icc -Ofast laplace2d.c -o L2Di
```


Executing & Measuring Performance

Command to
instrument execution

Performance metrics
to measure

iter_max declared
at run-time

```
$ perf stat -e cycles,instructions,cache-misses,task-clock ./L2Dg 100
```

```
Jacobi relaxation Calculation: 4096 x 4096 mesh, maximum of 100 iterations
```

```
10, 0.155006
20, 0.110024
30, 0.089901
40, 0.077374
50, 0.069623
60, 0.063285
70, 0.058825
80, 0.054945
90, 0.051831
100, 0.049208
```

Mesh size is declared
at compile-time

Error after 100 iterations

```
Performance counter stats for './L2Dg 100':
```

| | | | | |
|-----------------|-------------------|---|--------|-----------------|
| 289.757.998.767 | cycles | # | 3,367 | GHz |
| 28.377.296.648 | instructions | # | 0,10 | insns per cycle |
| 5.631.065.118 | cache-misses | # | 65,433 | M/sec |
| 86057,833003 | task-clock (msec) | # | 1,001 | CPUs utilized |

```
85,940821112 seconds time elapsed
```

Performance metrics
of execution

Processor Performance

$$\text{CPU Time} = \text{clock cycles} \times \text{clock cycle time}$$

$$\text{CPU Time} = \text{instructions} / (\text{IPC} \times \text{clock frequency})$$

$$\text{clock cycle time} = 1 / \text{clock frequency}$$

IPC: ratio of Instructions executed Per clock Cycle

cache miss: the data requested by the program do not reside in the internal memory of the processor chip and must be retrieved from main memory (DRAM). Cache misses degrade the IPC metric

Final Optimized Code

Code Motion

```
while ( err > tol*tol && iter < iter_max ) {
```

```
    err = 0.0f;
```

```
    if (iter%2 == 0)
```

Double Buffer

```
        for( j=1; j < n-1; j++ )
```

```
            for ( i=1; i < m-1; i++ ) {
```

Loop
Interchange

```
                Anew[j][i]= (A[j][i+1]+A[j][i-1]+  
                            A[j-1][i]+A[j+1][i])*0.25f;
```

Strength
Reduction

Loop
Fusion

```
                err = fmaxf( err, fabsf(Anew[j][i]-A[j][i]));
```

```
            }
```

```
    else
```

```
        for( j=1; j < n-1; j++ )
```

```
            for ( i=1; i < m-1; i++ ) {
```

```
                A[j][i]= (Anew[j][i+1]+Anew[j][i-1]+  
                            Anew[j-1][i]+Anew[j+1][i])*0.25f;
```

```
                err = fmaxf( err, fabsf(Anew[j][i]-A[j][i]));
```

```
            }
```

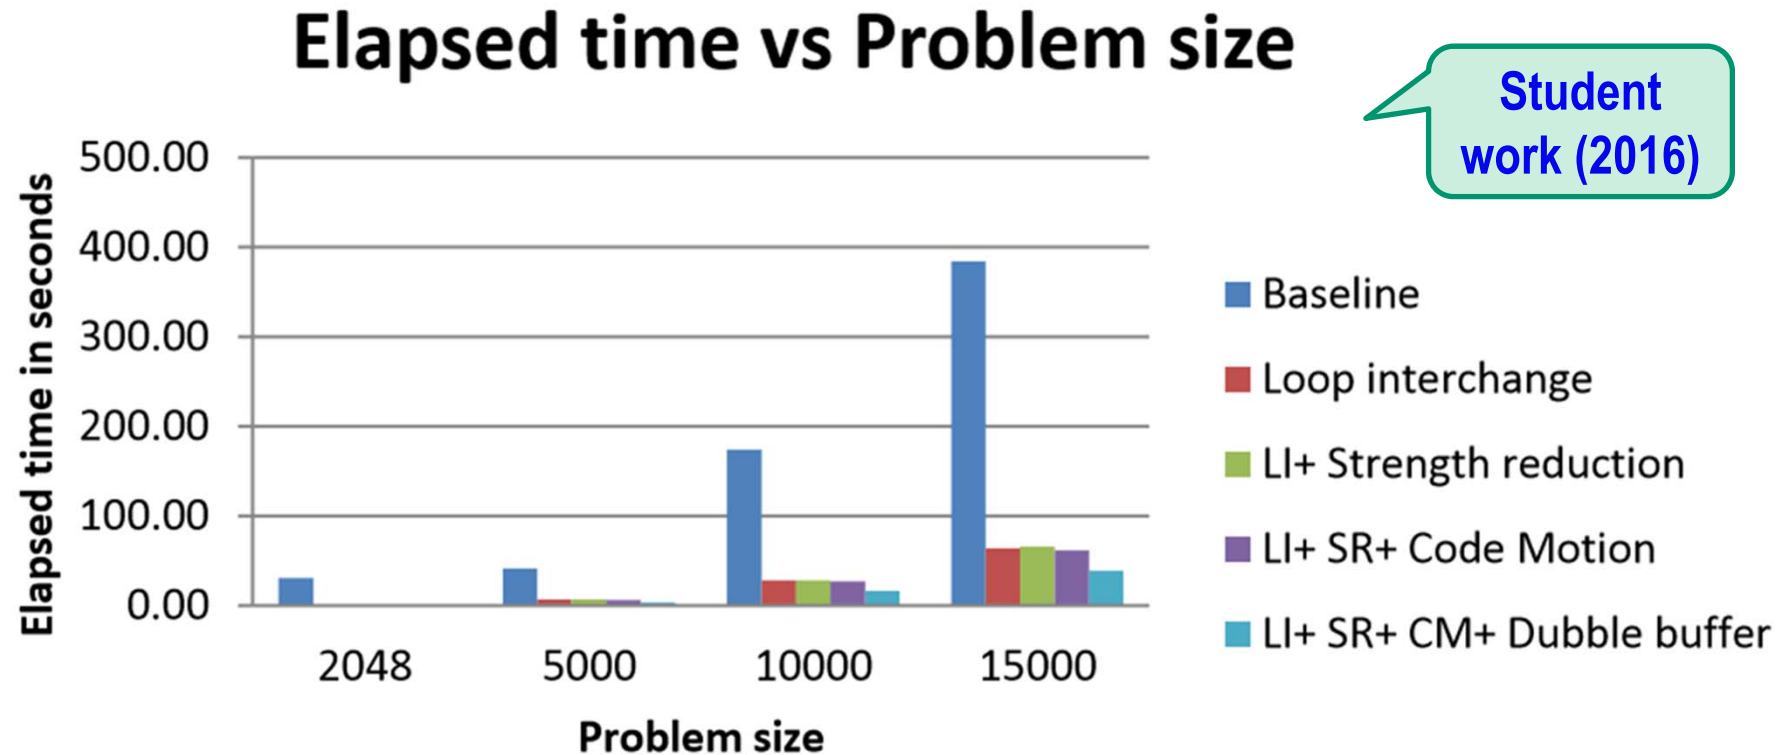
```
    if (++iter % (iter_max/10) == 0)
```

```
        printf("%5d, %0.6f\n", iter, sqrtf(err));
```

Code Motion

```
}
```

How to visualize (analyze) Performance



Elapsed time is not a good metric to compare variations on problem size

$$\text{Performance} = \frac{\text{Total Work}}{\text{Second}}$$

Better Metric
Normalized by Work

Work per second: Higher is Better

$$\frac{\text{Work}}{\text{Second}} = \frac{\text{Operations}}{\text{Instruction}} \frac{\text{Instructions}}{\text{Clock Cycle}} \frac{\text{Clock Cycles}}{\text{Second}}$$

OpRate **IPC** **Clock Freq.**

Coding Efficiency:
the more work of the
program done by a fixed
number of processor's
machine instructions,
the better


**MicroArchitecture
Throughput:**
The more machine
instructions executed every
clock cycle, the better

H/W speed:
The higher the clock
frequency of the
processor, the better

All factors are better when they are higher
The influence of the problem size is avoided
Programmers can improve OpRate & IPC

Algorithmic Efficiency: Complexity

```
Anew[j][i] = (A[j][i+1] + A[j][i-1] +  
              A[j-1][i] + A[j+1][i]) * 0.25f;  
error = fmaxf( error, fabsf(Anew[j][i] - A[j][i]));
```



5-point stencil

Constant WORK: on each iteration of inner loop


Stencil operations: $iter \times n \times m$

Algorithmic Complexity or Total Work:

$\theta (iter \times n \times m)$

Comparing Performance: Speedup

Problem Size = $2K \times 2K$ matrix, 100 convergence iterations



Loop interchange:

very large improvement (speedup is 21.4x)

Double Buffer:

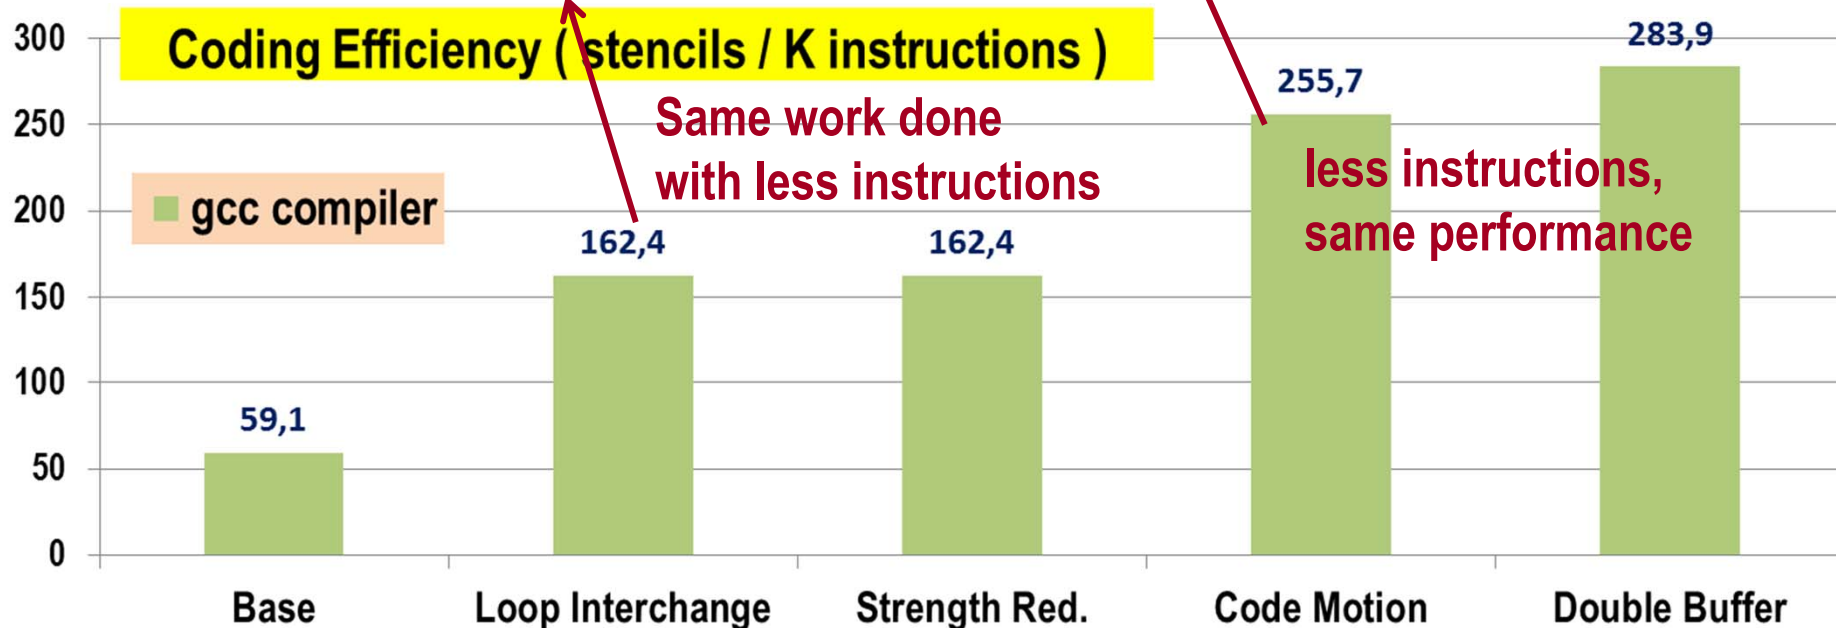

almost doubles performance (1.92x)

Strength Red. & Code motion:


very small or no improvement

Speedup:
 $\text{Perf}_{\text{NEW}} / \text{Perf}_{\text{OLD}}$

Explaining Speedup: Coding Efficiency



Explain: Processor Throughput



Profile & Assembly Inspection

```
$ module add gcc/6.1
$ gcc -Ofast -lm Laplace2D.c -o L2Dg
$ perf record ./L2Dg
$ perf report
```

generates file perf.data

uses file perf.data

| | | | | |
|--------|-------|-------------------|-----|----------------------|
| 99,70% | saxpy | L2Dg | [.] | main |
| 0,10% | saxpy | [kernel.kallsyms] | [k] | clear_page_c |
| 0,03% | saxpy | [kernel.kallsyms] | [k] | apic_timer_interrupt |
| 0,02% | saxpy | [kernel.kallsyms] | [k] | task_tick_fair |
| 0,01% | ... | | | |

Press enter here to
visualize annotated code

Assembly Code: analysis & comparison

```
0,23 1c8: movss (%rsi), %xmm4
1,17   add    $0x2, %r11d
      addss  (%rcx), %xmm4
1,22   movss (%rdx), %xmm5
42,47  addss  %xmm5, %xmm0
1,31   movss 0x4000(%rdx), %xmm7
22,47  add    $0x8000, %rsi
0,03   add    $0x8000, %rcx
0,14   add    $0x8000, %rdx
0,03   add    $0x8000, %rax
0,26   addss  %xmm0, %xmm4
0,17   movss  -0x4000(%rcx), %xmm0
0,94   addss  -0x4000(%rsi), %xmm0
1,08   mulss  %xmm3, %xmm4
0,43   movss  %xmm4, -0x8000(%rax)
0,09   subss  %xmm6, %xmm4
0,06   addss  %xmm7, %xmm6
0,60   andps  %xmm2, %xmm4
0,06   addss  %xmm6, %xmm0
0,43   movaps %xmm7, %xmm6
      sqrtss %xmm4, %xmm4
0,66   mulss  %xmm3, %xmm0
0,97   movss  %xmm0, -0x4000(%rax)
0,17   subss  %xmm5, %xmm0
0,51   cmp    $0xffd, %r11d
0,03   andps  %xmm2, %xmm0
0,34   sqrtss %xmm0, %xmm0
1,99   maxss  %xmm4, %xmm0
0,60   maxss  %xmm0, %xmm1
0,74   movaps %xmm5, %xmm0
jne 1c8
```

gcc -Ofast (v 6.1)

```
7,04 290: movups (%r9,%rax,1), %xmm0
0,19   movaps %xmm6, %xmm10
      movups (%rdx,%rax,1), %xmm9
1,97   addps  %xmm9, %xmm0
5,82   addps  (%r8,%rax,1), %xmm0
0,94   addps  (%rdi,%rax,1), %xmm0
11,73  mulps  %xmm5, %xmm0
0,47   movaps %xmm0, (%rcx,%rax,1)
0,28   subps  (%rsi,%rax,1), %xmm0
6,85   add    $0x10, %rax
0,09   cmp    $0x3fe0, %rax
      andps  %xmm4, %xmm0
6,19   cmpneq %xmm0, %xmm10
0,75   rsqrt  %xmm0, %xmm9
1,03   andps  %xmm10, %xmm9
0,56   mulps  %xmm9, %xmm0
7,97   mulps  %xmm0, %xmm9
5,16   mulps  %xmm2, %xmm0
2,44   addps  %xmm3, %xmm9
6,10   mulps  %xmm9, %xmm0
20,26  maxps  %xmm0, %xmm1
13,23  jne 290
```

Loop Fusion

Loop Interchange

mulps: Multiply Packed Single-Precision Floating-Point Values

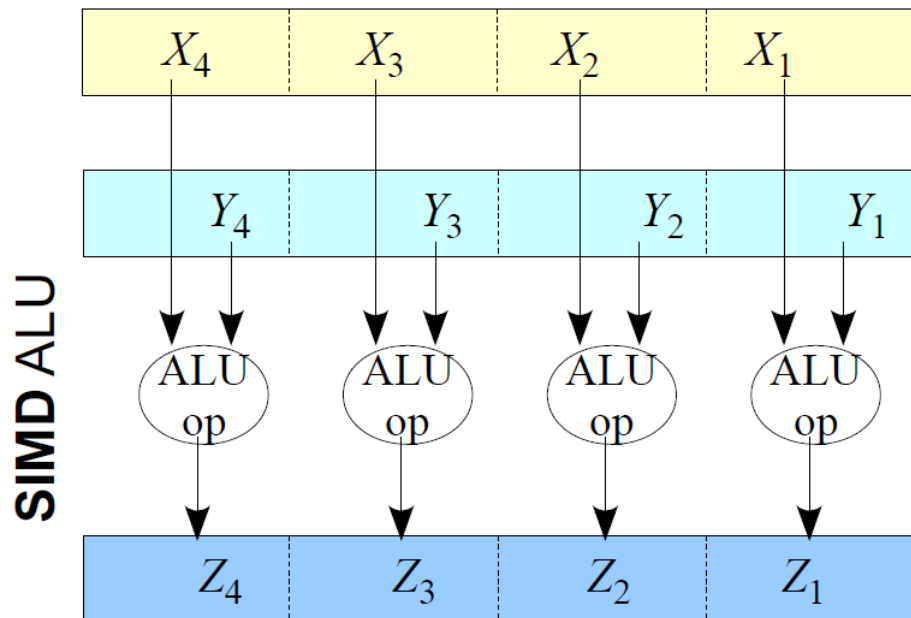
| Opcode | Mnemonic | Description |
|----------|-----------------------|---|
| 0F 59 /r | MULPS xmm1, xmm2/m128 | Multiply packed single-precision floating-point values in xmm2/mem by xmm1. |

Description

Performs an SIMD multiply of the four packed single-precision floating-point values from the source operand (second operand) and the destination operand (first operand), and stores the packed single-precision floating-point results in the destination operand. The source operand can be an XMM register or a 128-bit memory location. The destination operand is an XMM register. See Figure 10-5 in the IA-32 Intel Architecture Software Developer's Manual, Volume 1 for an illustration of an SIMD single-precision floating-point operation.

Operation

```
Destination[0..31] = Destination[0..31] * Source[0..31];  
Destination[32..63] = Destination[32..63] * Source[32..63];  
Destination[64..95] = Destination[64..95] * Source[64..95];  
Destination[96..127] = Destination[96..127] * Source[96..127];
```



4 pairs of input data operands

4 SIMD results produced simultaneously.
SIMD: Single-Instruction Multiple-Data

mulss: Multiply Scalar Single-Precision Floating-Point Values


| Opcode | Mnemonic | Description |
|-------------|----------------------|--|
| F3 0F 59 /r | MULSS xmm1, xmm2/m32 | Multiply the low single-precision floating-point value in xmm2/mem by the low single-precision floating-point value in xmm1. |

Description

Multiplies the low single-precision floating-point value from the source operand (second operand) by the low single-precision floating-point value in the destination operand (first operand), and stores the single-precision floating-point result in the destination operand. The source operand can be an XMM register or a 32-bit memory location. The destination operand is an XMM register. The three high-order doublewords of the destination operand remain unchanged. See Figure 10-6 in the IA-32 Intel Architecture Software Developer's Manual, Volume 1 for an illustration of a scalar single-precision floating-point operation.

Operation

```
Destination[0..31] = Destination[0..31] * Source[0..31];  
//Destination[32..127] remains unchanged
```




1 pair of input data operands.
Only the less significant 32 bits are used

1 scalar result.
Remaining 96 bits unmodified


Example: Scalar versus Vector Addition

Scalar Addition



Vector Addition

Vector Lanes (8 in this example)

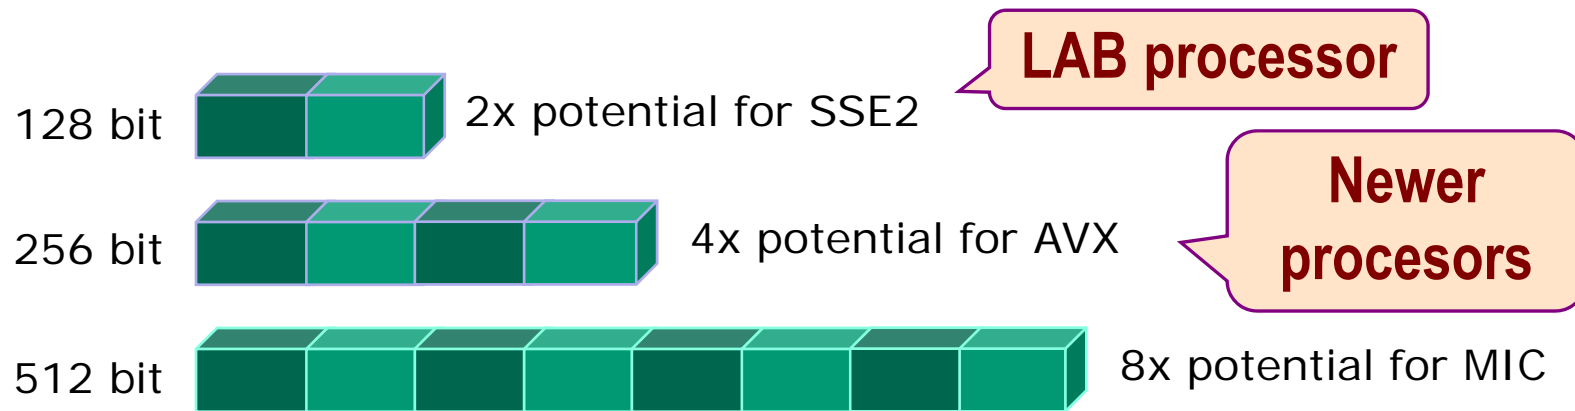


Vector length [in number of elements] =
size of vector register [in bits] / size of the data type [in bits]

Number of **vector lanes** = Vector Length

Potential Performance Speedups

Double Precision FP vector width vs
theoretical speedup potential over scalar



128, 256, 512 bit vector divided by **64 bit data type** yields
potential speedups of **2, 4, or 8** times

- **Wider vectors** allow for **higher potential performance** gains
- Gains of 4X and 8X within reach using vectorization capability

Ways to Write Vector Code C/C++

Data Level Parallelism with OpenMP* 4.0

Serial Code

```
for(i = 0; i < N; i++)  
    A[i] = B[i] + C[i];
```

Let compiler decide if vectorizing is safe

SIMD Pragma/Directive

```
#pragma omp simd  
for(i = 0; i < N; i++)  
    A[i] = B[i] + C[i];
```

Tell compiler it is safe to vectorize

SIMD-enabled Function

```
#pragma omp declare simd  
float foo(float B, float C)  
    return B + C;
```

...

```
// call foo below  
#pragma omp simd  
for(i = 0; i < N; i++)  
    A[i] = foo(B[i], C[i]);
```

A function inside a vectorized loop must have a vectorized version

```
$ icc -Ofast -openmp file.c -o ...  
$ gcc -Ofast -fopenmp file.c -o ...
```


Requirements for Loop Vectorization

- **Independent iterations:**

Several loop iterations can be executed in parallel

```
for (i=1; i<N; i++)  
    A[i] = A[i-1] + B[i];
```

Not vectorizable:
Iteration i depends
on data produced
by iteration $i-1$

Requirements for Loop Vectorization

- **Independent iterations:**

Several loop iterations can be executed in parallel

- **Consecutive Data Accesses:**

Consecutive loop iterations access consecutive memory locations

```
for (i=0; i<N; i++)  
    A[i][5] = A[i][5] + B[i][2];
```

READ & WRITE DATA no vectorizable:

**Cannot read elements A[i][5] and A[i+1][5] from memory
(or B[i][2] and B[i+1][2])
with a single SIMD Load instruction**

Requirements for Loop Vectorization

- **Independent iterations:**

Several loop iterations can be executed in parallel

- **Consecutive Data Accesses:**

Consecutive loop iterations access consecutive memory locations

- **No conditional divergence:**

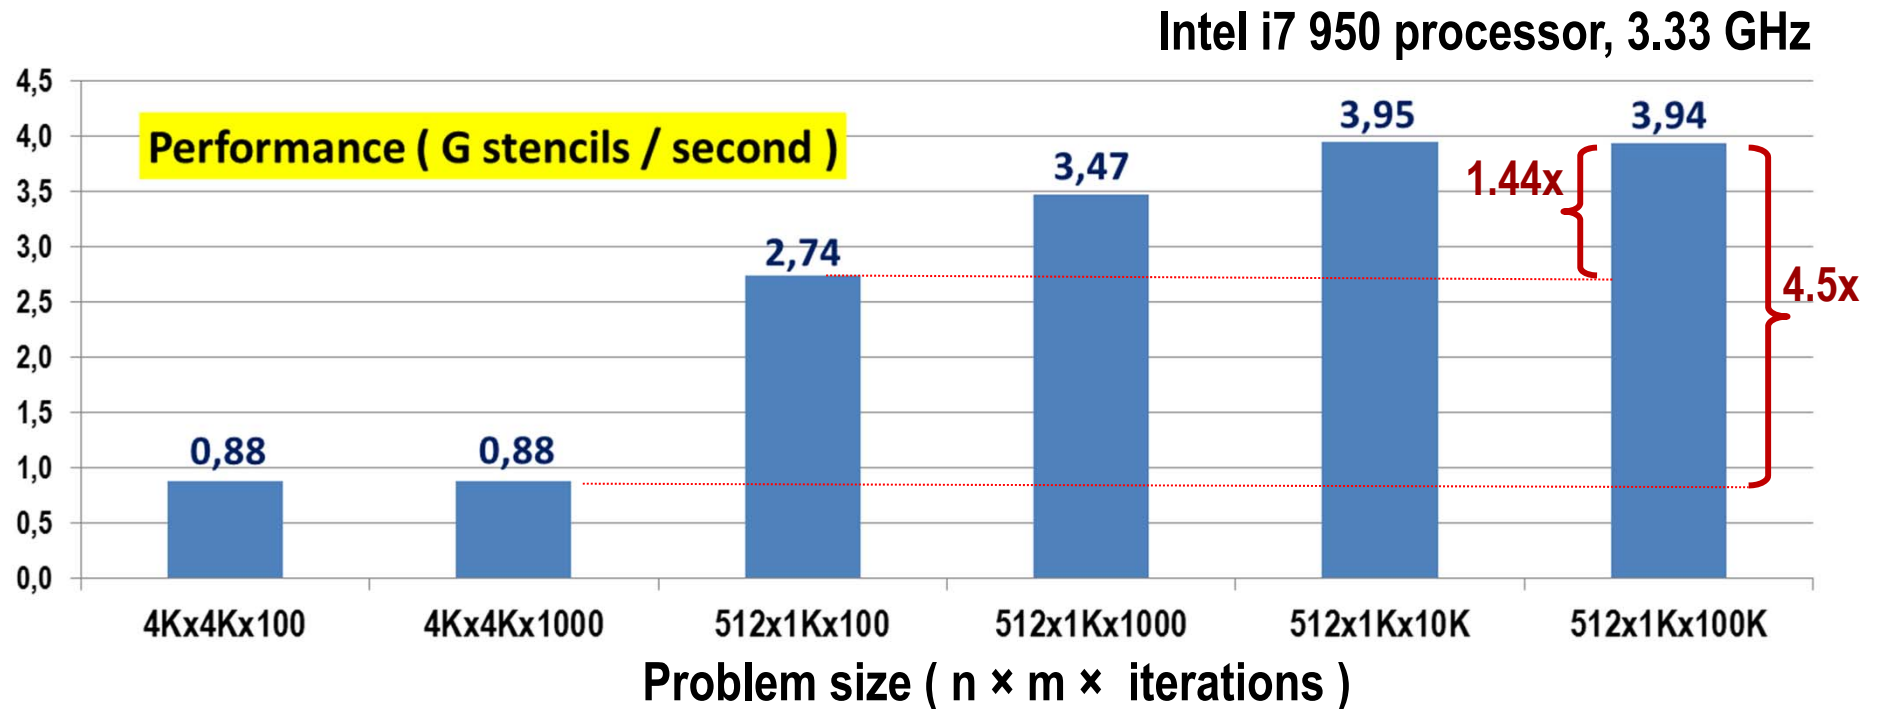
Conditional statements: some computation done depending on the condition outcome

```
for (i=0; i<N; i++)  
    if (A[i]>5) B[i] = A[i]-B[i];
```

Not efficient vectorization:

Masked execution: only some SIMD lanes are active


Effect of Problem Size



Work / second metric simplifies comparisons for different problem sizes.


How to explain the results?

Explained by Coding Efficiency?




Intel i7 950 processor,
3.33 GHz

Not explained by
looking at the number
of machine
instructions executed



Microarchitecture Throughput?




Intel i7 950 processor,
3.33 GHz


Still cannot
explain this ...

Probably: printf() Input/Output

DO explain this
improvement



The CPU-DRAM Gap: the **Memory Wall**




During the latency of a DRAM read operation the processor can execute thousands of instructions

Memory Technology

Small memory is faster than Large memory

SRAM: very fast memory technology (inside die)

**DRAM: (off-chip)
more area-efficient &
energy-efficient
memory technology
(but much slower)**




Basic Idea of a Cache

Cache contains copies of some memory locations

Memory access does *cache lookup* first (L1 = Level 1).

If desired data is not in the cache, then **read** the data from main memory ...

and **store** data in cache (**replace** some previous data)



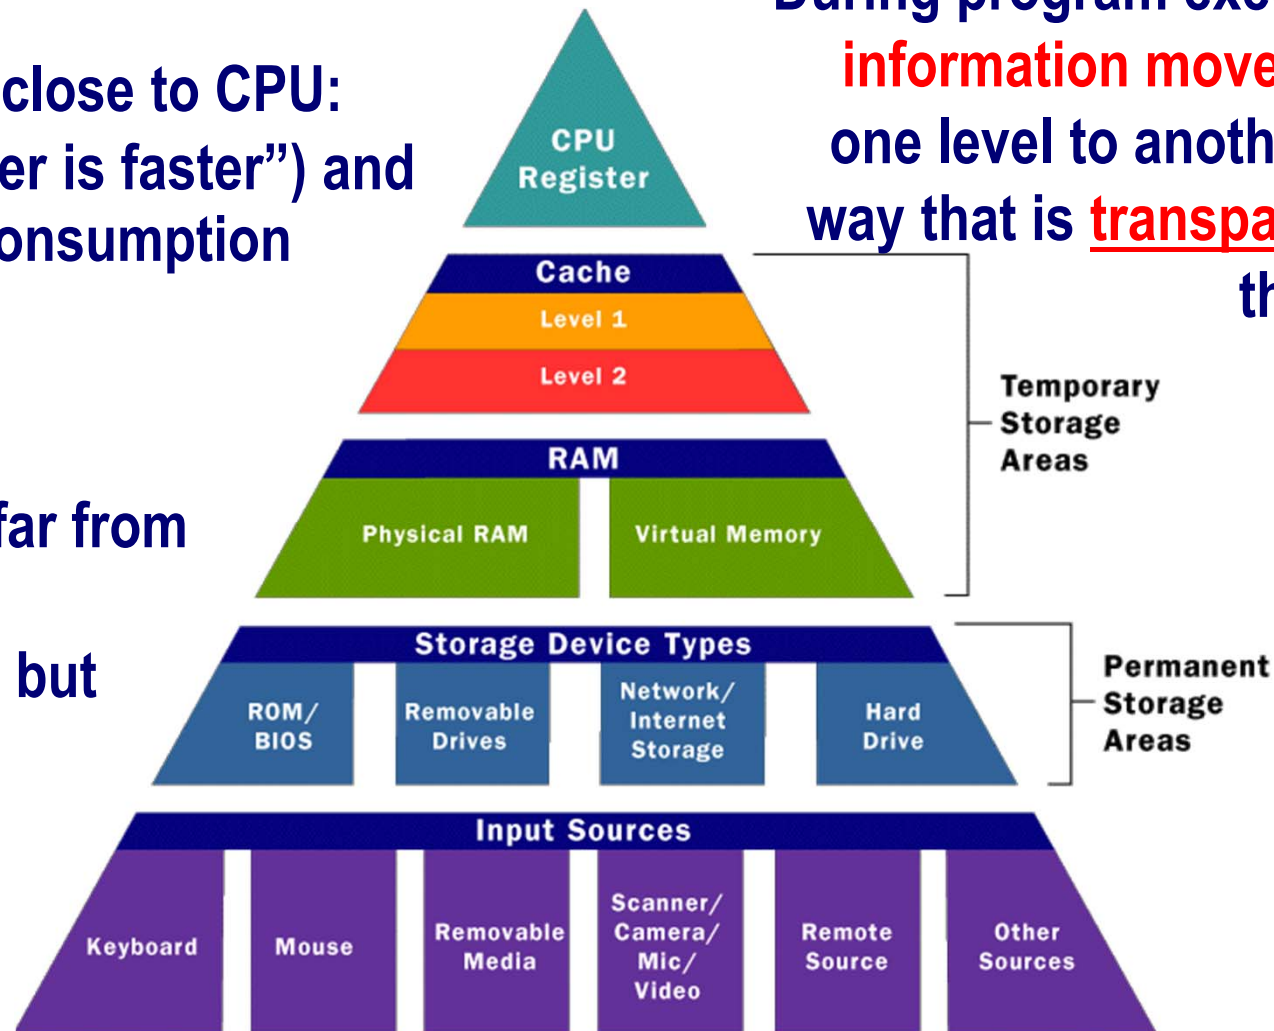
Memory Hierarchy

Illusion of fast, large, cheap, non-volatile, low-energy memory

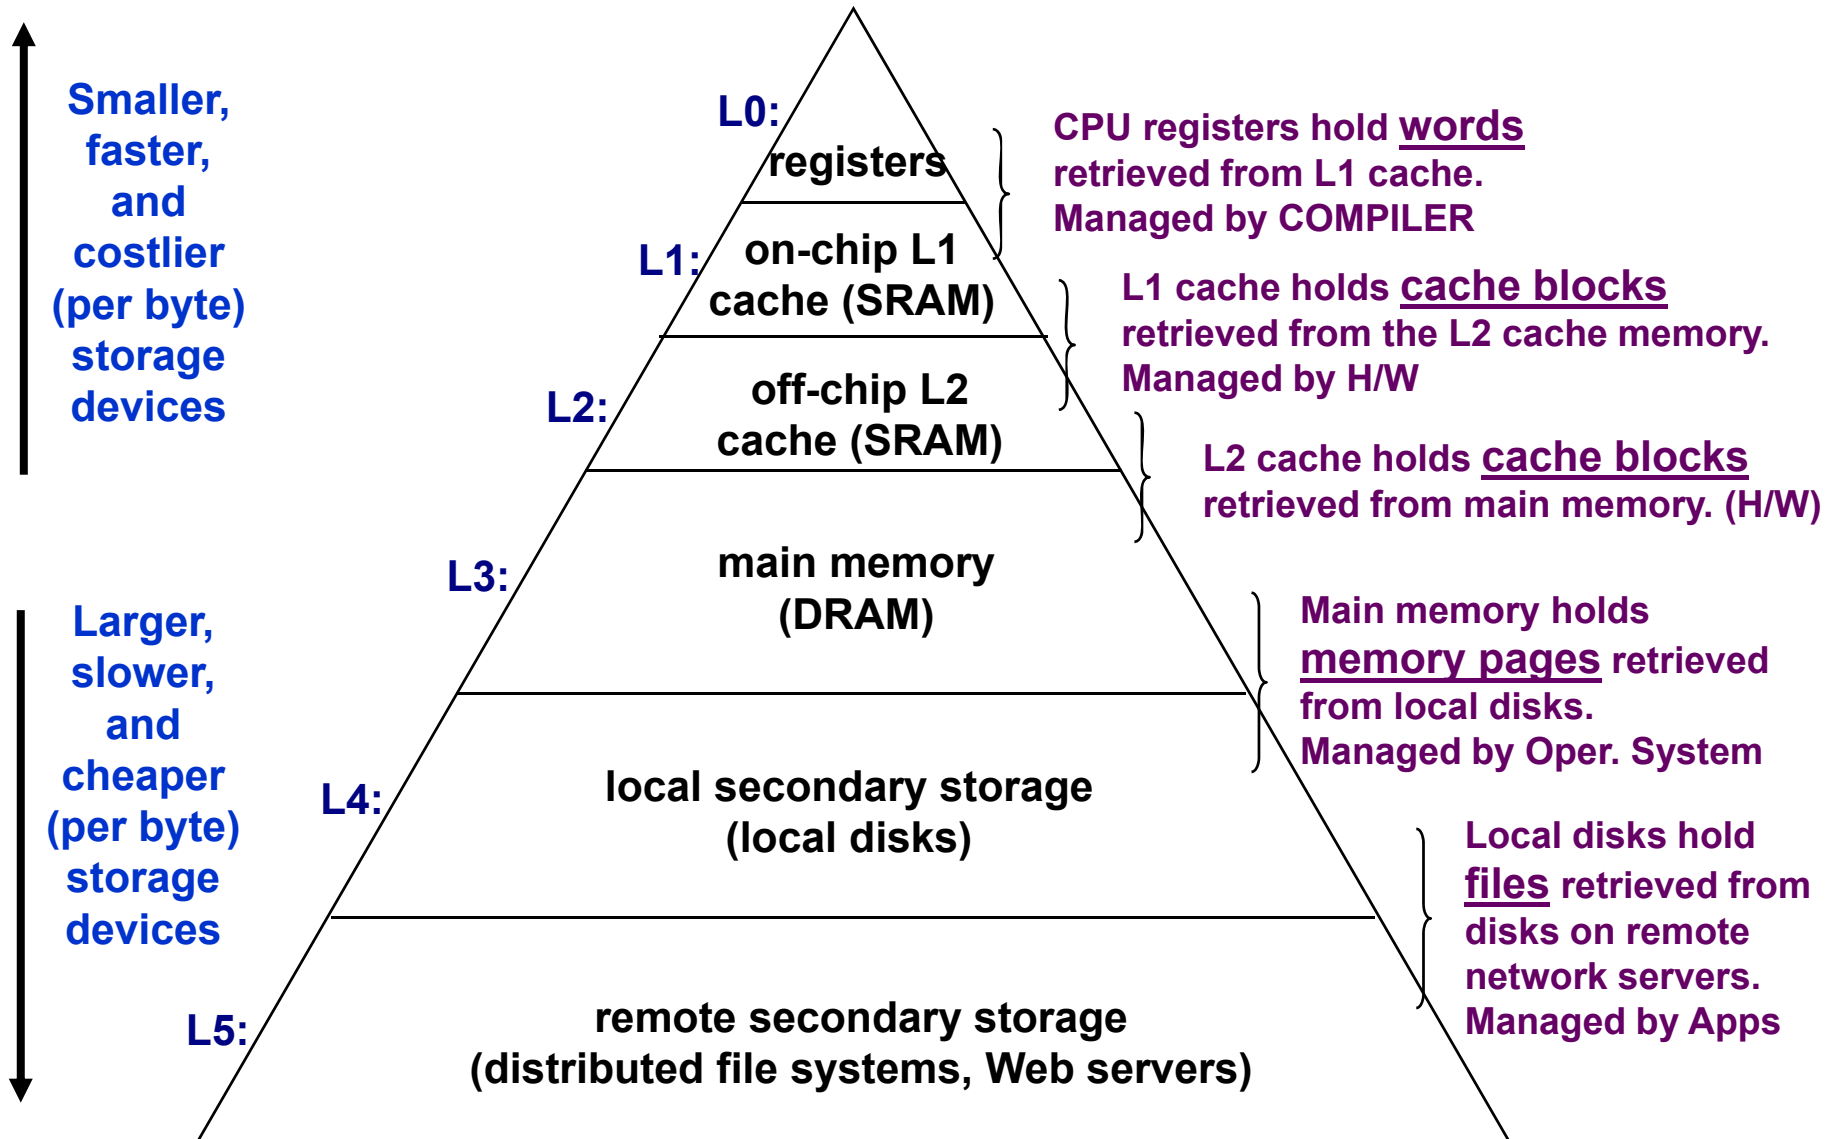
Memory levels close to CPU:
faster. (“smaller is faster”) and
lower energy consumption

Memory levels far from
CPU: **cheaper**
(larger capacity but
slower)

During program execution,
information moves from
one level to another, in a
way that is transparent to
the user




Memory Hierarchy (2)

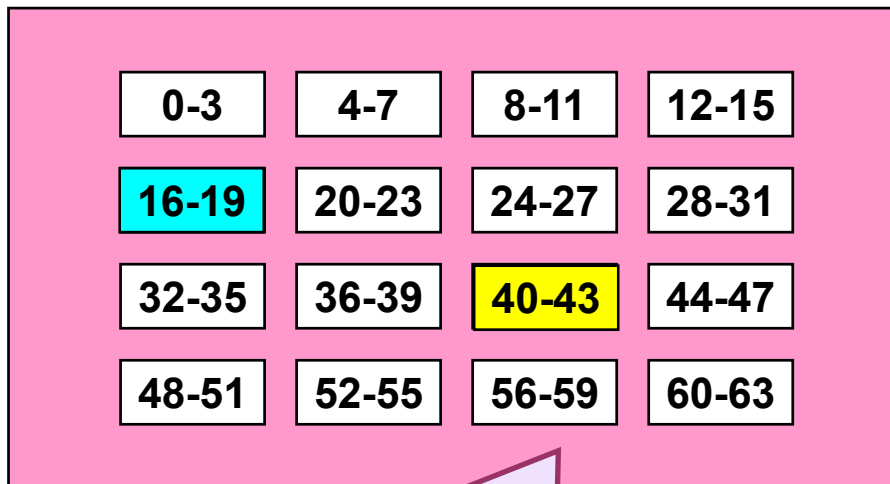


Cache Organization in Blocks

Cache:



Memory:



Data copied in
block-sized
transfer units

40-43

Blocks are aligned into memory

Smaller, faster, more expensive
memory device
(caches a subset of the
memory blocks into cache lines)

Cache line: physical
storage for a block of
memory values

Larger, slower, cheaper storage
device is partitioned
into memory blocks

Memory block: logical
partition of consecutive
memory locations

Cache Organization in Blocks (2)

Ca

Why Blocks?
Exploit Spatial Locality

Smaller, faster, more expensive
memory device
(caches a subset of the
memory blocks into cache lines)

Data copied in

Memory



Rule of Thumb for PERFORMANCE:

When the program requests a new data block from memory, most of the data in the block should be used (maybe several times) in a relatively short period of time

This naturally occurs when data is read from consecutive positions in an array


Blocks are

Measure Memory Cache Performance



Comparing Performance: Speedup

Problem Size = $2K \times 2K$ matrix, 100 convergence iterations



Loop interchange:

very large improvement (speedup is 21.4x)

Double Buffer:

almost doubles performance (1.92x)

Strength Red. & Code motion:

very small or no improvement

Assembly: analysis & comparison (2)

gcc -Ofast (v 6.1)

| | | | | | |
|-------|------|-----------------------------|-------|------|-----------------------------|
| 7,04 | 290: | movups (%r9,%rax,1),%xmm0 | 7,75 | 290: | movups (%r9,%rax,1),%xmm0 |
| 0,19 | | movaps %xmm6,%xmm10 | 0,96 | | movaps %xmm6,%xmm10 |
| | | movups (%rdx,%rax,1),%xmm9 | | | movups (%rdx,%rax,1),%xmm9 |
| 1,97 | | addps %xmm9,%xmm0 | 1,44 | | addps %xmm9,%xmm0 |
| 5,82 | | addps (%r8,%rax,1),%xmm0 | 5,26 | | addps (%r8,%rax,1),%xmm0 |
| 0,94 | | addps (%rdi,%rax,1),%xmm0 | | | addps (%rdi,%rax,1),%xmm0 |
| 11,73 | | mulps %xmm5,%xmm0 | | | mulps %xmm5,%xmm0 |
| 0,47 | | movaps %xmm0, (%rcx,%rax,1) | | | movaps %xmm0, (%rcx,%rax,1) |
| 0,28 | | subps (%rsi,%rax,1),%xmm0 | | | subps (%rsi,%rax,1),%xmm0 |
| 6,85 | | add \$0x10,%rax | | | add \$0x10,%rax |
| 0,09 | | cmp \$0x3fe0,%rax | | | cmp \$0x3fe0,%rax |
| | | andps %xmm4,%xmm0 | | | andps %xmm4,%xmm0 |
| 6,19 | | cmpneq %xmm0,%xmm10 | 0,10 | | andps %xmm4,%xmm0 |
| 0,75 | | rsqrtp %xmm0,%xmm9 | 4,88 | | cmpneq %xmm0,%xmm10 |
| 1,03 | | andps %xmm10,%xmm9 | 0,38 | | rsqrtp %xmm0,%xmm9 |
| 0,56 | | mulps %xmm9,%xmm0 | 0,67 | | andps %xmm10,%xmm9 |
| 7,97 | | mulps %xmm0,%xmm9 | 0,29 | | mulps %xmm9,%xmm0 |
| 5,16 | | mulps %xmm2,%xmm0 | 8,42 | | mulps %xmm0,%xmm9 |
| 2,44 | | addps %xmm3,%xmm9 | 6,51 | | mulps %xmm2,%xmm0 |
| 6,10 | | mulps %xmm9,%xmm0 | 2,97 | | addps %xmm3,%xmm9 |
| 20,26 | | maxps %xmm0,%xmm1 | 6,51 | | mulps %xmm9,%xmm0 |
| 13,23 | | jne 290 | 17,80 | | maxps %xmm0,%xmm1 |
| | | | 16,56 | | jne 290 |

Same Exact CODE:
Compiler already performs
this optimization

Loop interchange

Strength reduction

Assembly: analysis & comparison (3)

```

7,75 290: movups  (%r9,%rax,1),%xmm0
0,96    movaps  %xmm6,%xmm10
      movups  (%rdx,%rax,1),%xmm9
1,44    addps   %xmm9,%xmm0
5,26    addps  (%r8,%rax,1),%xmm0
0,67    addps  (%rdi,%rax,1),%xmm0
12,54   mulps  %xmm5,%xmm0
0,29    movaps  %xmm0, (%rcx,%rax,1)
0,10    subps  (%rsi,%rax,1),%xmm0
4,98    add     $0x10,%rax
0,10    cmp     $0x3fe0,%rax
4,88    andps   %xmm4,%xmm0
0,38    cmpeq   %xmm0,%xmm10
0,67    andps   %xmm10,%xmm9
0,29    mulps   %xmm9,%xmm0
8,42    mulps   %xmm0,%xmm9
6,51    mulps   %xmm2,%xmm0
2,97    addps   %xmm3,%xmm9
6,51    mulps   %xmm9,%xmm0
17,80   maxps   %xmm0,%xmm1
16,56   jne     290
    
```

Code Motion:
remove
computation of
square root from
inner loop


Strength reduction

```

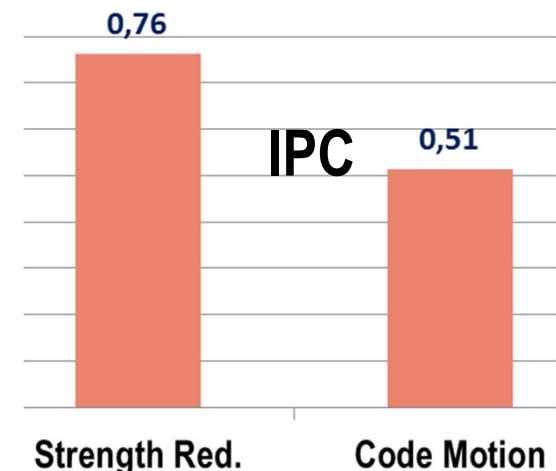
7,53 278: movups  (%r9,%rax,1),%xmm0
1,06    movups  (%rdx,%rax,1),%xmm2
6,71    addps   %xmm2,%xmm0
1,18    addps  (%r8,%rax,1),%xmm0
9,65    addps  (%rdi,%rax,1),%xmm0
33,29   mulps  %xmm4,%xmm0
12,00   movaps  %xmm0, (%rcx,%rax,1)
2,12    subps  (%rsi,%rax,1),%xmm0
9,53    add     $0x10,%rax
2,82    cmp     $0x3fe0,%rax
13,65   andps   %xmm3,%xmm0
      maxps   %xmm0,%xmm1
      jne     278
    
```

gcc -Ofast (v 6.1)

Performance bottleneck is waiting for data in memory: reducing instructions does not improve execution time.



Why IPC degrades?



Assembly: analysis & comparison (3)

Samples: 154K of event 'cycles', Event count (approx.): 128527894268

| Overhead | Command | Shared Object | Symbol |
|----------|---------|-------------------|--------------------------|
| 53,27% | Lsqr | Lsqr | [.] main |
| 45,82% | Lsqr | libc-2.17.so | [.] __memcpy_ssse3_back |
| 0,04% | Lsqr | [kernel.kallsyms] | [k] update_wall_time |
| 0,04% | Lsqr | [kernel.kallsyms] | [k] apic_timer_interrupt |
| 0,04% | Lsqr | [kernel.kallsyms] | [k] rcu_check_callbacks |
| 0,03% | Lsqr | [kernel.kallsyms] | [k] trigger_load_balance |

Code Motion

gcc -Ofast (v 6.1)

Double Buffer

Samples: 94K of event 'cycles', Event count (approx.): 76992055233

| Overhead | Command | Shared Object | Symbol |
|----------|---------|-------------------|---------------|
| 99,21% | L2buff | L2buff | [.] main |
| 0,05% | L2buff | [kernel.kallsyms] | [k] apic_time |
| 0,03% | L2buff | [kernel.kallsyms] | [k] rcu_check |
| 0,03% | L2buff | [kernel.kallsyms] | [k] clear_pag |
| 0,03% | L2buff | [kernel.kallsyms] | [k] task_tick |
| 0,02% | L2buff | [kernel.kallsyms] | [k] trigger_1 |
| 0,02% | L2buff | libc-2.17.so | [.] __memset |

The few instructions doing the memory copy were responsible of most of the cycles

Why IPC improves?

