

Shared Memory - OpenMP

Dept. Arquitectura de Computadors i Sistemes Operatius
Universitat Autònoma de Barcelona



**Universitat
Autònoma
de Barcelona**

Shared Memory

- › **OpenMP is currently the programming standard for the share memory model on multicore systems.**

Shared Memory

- › Thread based model.
- › Threads read and write shared variables.
- › Synchronization mechanism are offered.
- › It is possible to change the attributes of threads and data for minimizing synchronization.

OpenMP

- › **OpenMP is not an automatic parallelization tool:**
 - Programmers must specify parallelism explicitly.
- › **OpenMP is not only for exploiting loop parallelism:**
 - It also offers functionalities for other forms of parallelism.

OpenMP

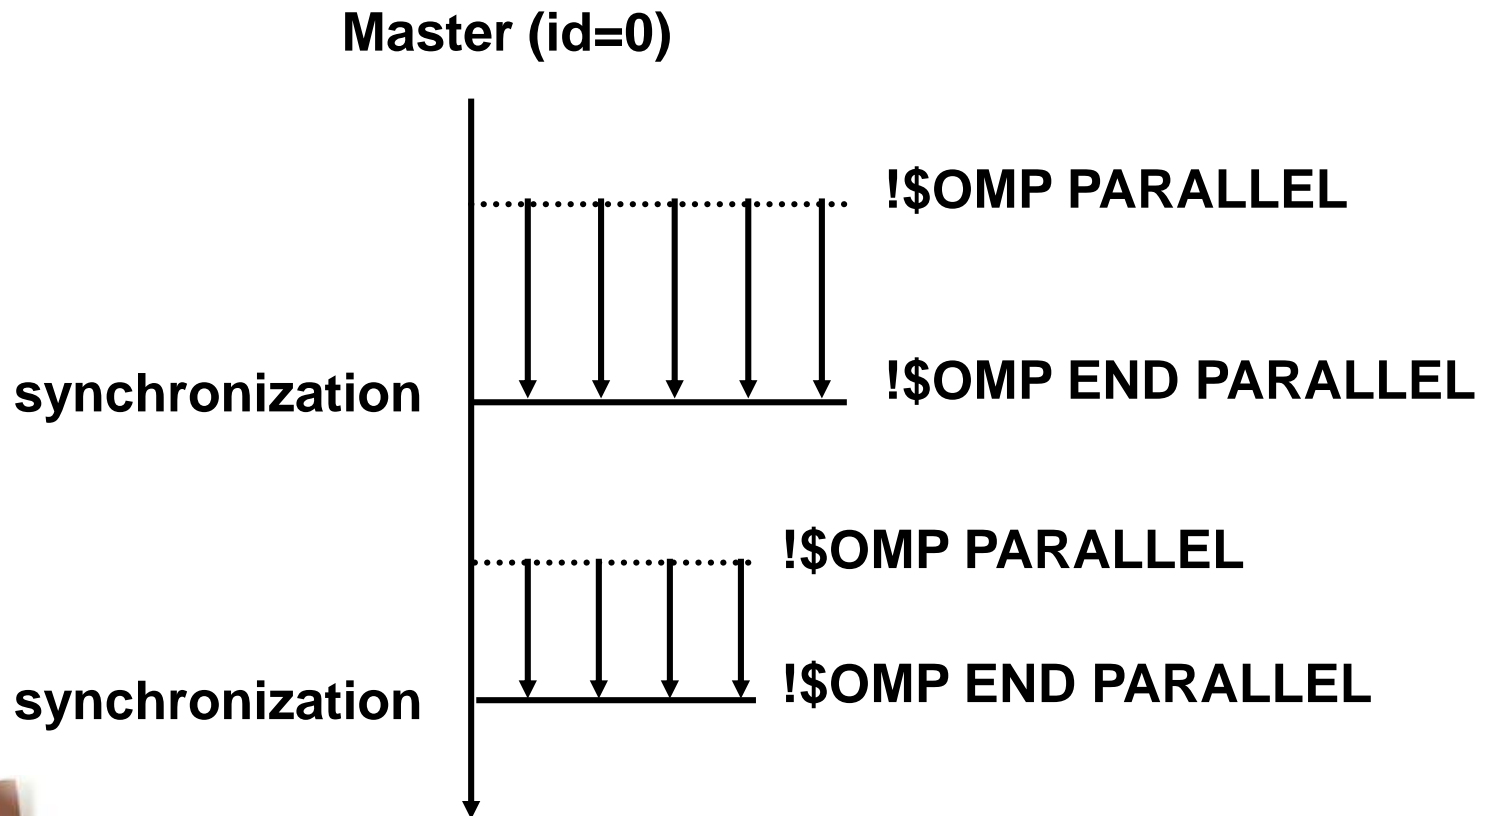
- › **OpenMP is not a programming language:**
 - It is structured as extensions using directives to base languages like Fortran or C.
- › **OpenMP is not only a research project:**
 - Many commercial compilers support OpenMP.

OpenMP

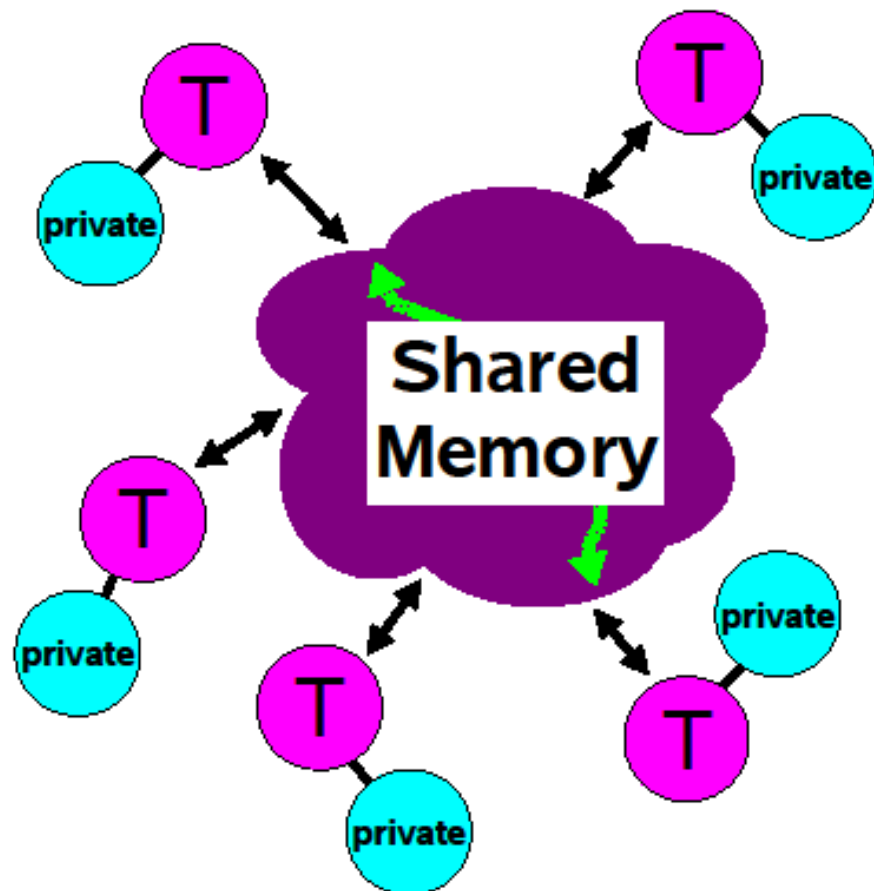
- › **OpenMP is an API (Application Program Interface) for parallel programming in shared memory systems.**
- › **It's main objective is to easily parallelize existing applications.**
- › **It's based in directives introduced in the program as special comments (pragmas).**

OpenMP Execution Model

- › Initially based on a FORK-JOIN model



Shared Memory Model



Programming Model

- ✓ All threads have access to the same, globally shared, memory
- ✓ Data can be shared or private
- ✓ Shared data is accessible by all threads
- ✓ Private data can be accessed only by the threads that owns it
- ✓ Data transfer is transparent to the programmer
- ✓ Synchronization takes place, but it is mostly implicit

About Data



◆ *In a shared memory parallel program variables have a "label" attached to them:*

☞ *Labelled "Private" ➤ Visible to one thread only*

✓ *Change made in local data, is not seen by others*

✓ *Example - Local variables in a function that is executed in parallel*

☞ *Labelled "Shared" ➤ Visible to all threads*

✓ *Change made in global data, is seen by all others*

✓ *Example - Global data*

Components of OpenMP



Directives

- ◆ *Parallel regions*
- ◆ *Work sharing*
- ◆ *Synchronization*
- ◆ *Data scope attributes*
 - ☞ *private*
 - ☞ *firstprivate*
 - ☞ *lastprivate*
 - ☞ *shared*
 - ☞ *reduction*
- ◆ *Orphaning*

Environment variables

- ◆ *Number of threads*
- ◆ *Scheduling type*
- ◆ *Dynamic thread adjustment*
- ◆ *Nested parallelism*

Runtime environment

- ◆ *Number of threads*
- ◆ *Thread ID*
- ◆ *Dynamic thread adjustment*
- ◆ *Nested parallelism*
- ◆ *Timers*
- ◆ *API for locking*

User Interface

› Compiler directives

- There are control structures and data attributes structures.
- Compilers ignore these directives (they are just comments) unless the proper options are used when compiling (“-mp” or “-fopenmp”).

User Interface

› Library

- Set of functions for controlling some parameters, such as the number of threads to be

```
call omp_set_num_threads (128)
```

› But also environment variables

- Another way of doing the same

```
setenv OMP_NUM_THREADS 8
```

The parallel region



A parallel region is a block of code executed by multiple threads simultaneously

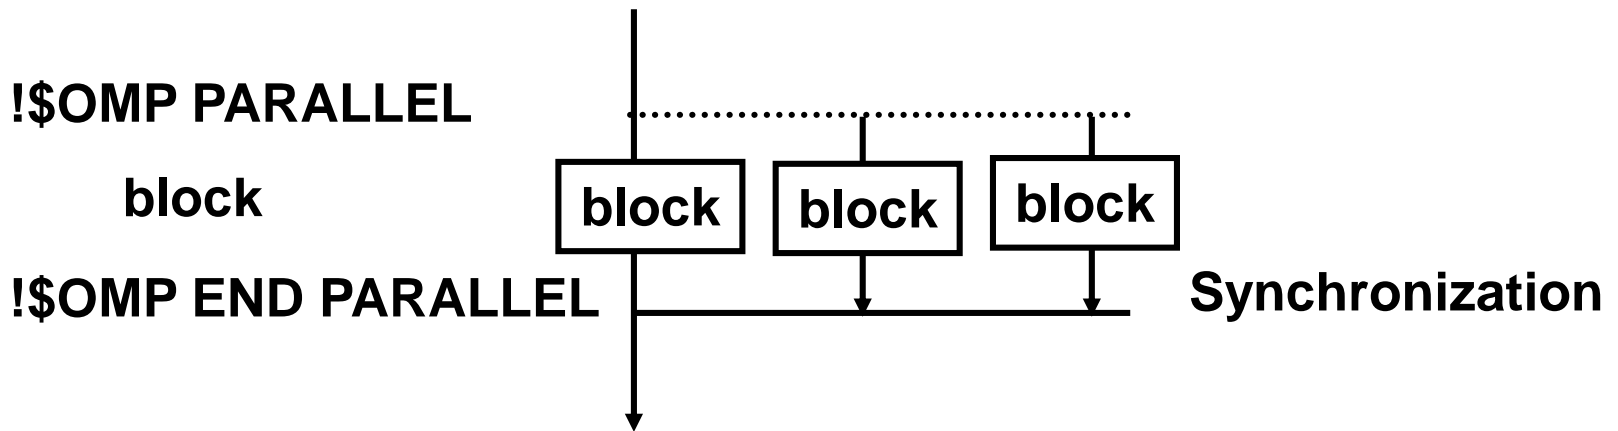
```
#pragma omp parallel [clause[,] clause] ...  
{  
    "this will be executed in parallel"  
} (implied barrier)
```

```
!$omp parallel [clause[,] clause] ...  
    "this will be executed in parallel"  
!$omp end parallel (implied barrier)
```

Directives

› PARALLEL / END PARALLEL

- Define a parallel region
- It does the “fork” and “join”.
- The number of threads is constant in the parallel regions



The parallel region - clauses



A parallel region supports the following clauses:

if	<i>(scalar expression)</i>	
private	<i>(list)</i>	
shared	<i>(list)</i>	
default	<i>(none shared)</i>	<i>(C/C++)</i>
default	<i>(none shared private)</i>	<i>(Fortran)</i>
reduction	<i>(operator: list)</i>	
copyin	<i>(list)</i>	
firstprivate	<i>(list)</i>	
num_threads	<i>(scalar_int_expr)</i>	

lastprivate(list)

A more elaborate example



```
#pragma omp parallel if (n>limit) default(none) \
    shared(n,a,b,c,x,y,z) private(f,i,scale)
{
```

```
    f = 1.0;
```

```
#pragma omp for nowait
```

```
    for (i=0; i<n; i++)
        z[i] = x[i] + y[i];
```

```
#pragma omp for nowait
```

```
    for (i=0; i<n; i++)
        a[i] = b[i] + c[i];
```

```
#pragma omp barrier
```

```
    ....
    scale = sum(a,0,n) + sum(z,0,n) + f;
    ....
```

```
} /*-- End of parallel region --*/
```

Statement is executed
by all threads

parallel loop
(work will be distributed)

parallel loop
(work will be distributed)

synchronization

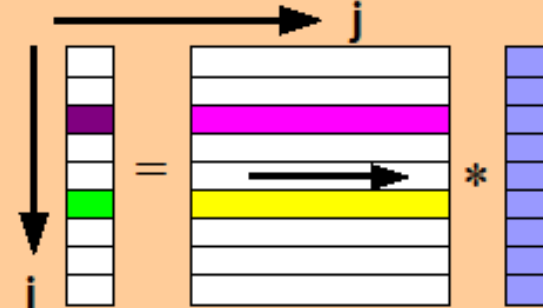
Statement is executed
by all threads

parallel region

Example - Matrix times vector



```
#pragma omp parallel for default(none) \
    private(i,j,sum) shared(m,n,a,b,c)
for (i=0; i<m; i++)
{
    sum = 0.0;
    for (j=0; j<n; j++)
        sum += b[i][j]*c[j];
    a[i] = sum;
}
```



TID = 0

for (i=0,1,2,3,4)

i = 0

sum = $\sum b[i=0][j]*c[j]$
a[0] = sum

i = 1

sum = $\sum b[i=1][j]*c[j]$
a[1] = sum

TID = 1

for (i=5,6,7,8,9)

i = 5

sum = $\sum b[i=5][j]*c[j]$
a[5] = sum

i = 6

sum = $\sum b[i=6][j]*c[j]$
a[6] = sum

... etc ...

The default clause



default (none | shared | private)

default (none | shared)

none

- ✓ *No implicit defaults*
- ✓ *Have to scope all variables explicitly*

shared

- ✓ *All variables are shared*
- ✓ *The default in absence of an explicit "default" clause*

private

- ✓ *All variables are private to the thread*
- ✓ *Includes common block data, unless **THREADPRIVATE***

Fortran

C/C++

Note: default(private) is not supported in C/C++

The reduction clause



reduction ([operator | intrinsic]) : list)

Fortran

reduction (operator : list)

C/C++

- ✓ *Reduction variable(s) must be shared variables*
- ✓ *A reduction is defined as:*

Fortran

```
x = x operator expr
x = expr operator x
x = intrinsic (x, expr_list)
x = intrinsic (expr_list, x)
```

C/C++

```
x = x operator expr
x = expr operator x
x++, ++x, x--, --x
x <binop> = expr
```

Check the docs
for details

- ✓ *Note that the value of a reduction variable is undefined from the moment the first thread reaches the clause till the operation has completed*
- ✓ *The reduction can be hidden in a function call*

The reduction clause - example



```
sum = 0.0
!$omp parallel default(none) &
!$omp shared(n,x) private(i)
!$omp do reduction (+:sum)
  do i = 1, n
    sum = sum + x(i)
  end do
!$omp end do
!$omp end parallel
print *,sum
```

Variable SUM is a shared variable

- ☞ *Care needs to be taken when updating shared variable SUM*
- ☞ *With the reduction clause, the OpenMP compiler generates code such that a race condition is avoided*

The nowait clause



- ❑ *To minimize synchronization, some OpenMP directives/pragmas support the optional **nowait** clause*
- ❑ *If present, threads will not synchronize/wait at the end of that particular construct*
- ❑ *In Fortran the nowait is appended at the closing part of the construct*
- ❑ *In C, it is one of the clauses on the pragma*

```
#pragma omp for nowait
{
    :
}
```

```
!$omp do
    :
    :
!$omp end do nowait
```

Directives

› Omp for - DO / END DO

- It's for classical parallel loops
- It must be in a parallel region
- Loop iterations are distributed among available threads
- Loop index is by default private to each thread

The omp for/do directive

The iterations of the loop are distributed over the threads

```
#pragma omp for [clause[[,] clause] ...]  
    <original for-loop>
```

```
!$omp do [clause[[,] clause] ...]  
    <original do-loop>  
!$omp end do [nowait]
```

Clauses supported:

private	firstprivate	
lastprivate	reduction	
<i>ordered*</i>	<i>schedule</i>	← <i>covered later</i>
nowait		

**) Required if ordered sections are in the dynamic extent of this construct*

DO / END DO

Program example

```
dimension A(100),B(100)
```

```
Integer i
```

```
!$OMP PARALLEL
```

```
!$OMP DO
```

```
    Do i=2,100
```

```
        B(i)= (A(i)+A(i-1))/2.0
```

```
    End Do
```

```
!$OMP END DO
```

```
!$OMP END PARALLEL
```

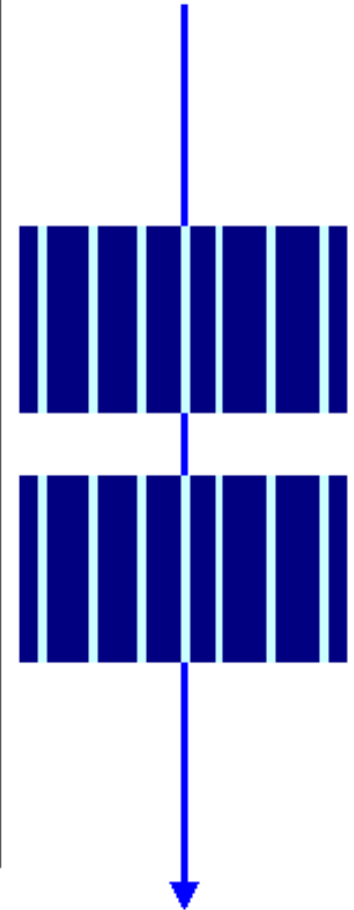
```
Return
```

```
End
```


The omp for directive - example



```
#pragma omp parallel default(none) \  
    shared(n,a,b,c,d) private(i)  
{  
    #pragma omp for nowait  
    for (i=0; i<n-1; i++)  
        b[i] = (a[i] + a[i+1])/2;  
  
    #pragma omp for nowait  
    for (i=0; i<n; i++)  
        d[i] = 1.0/c[i];  
  
} /*-- End of parallel region --*/  
    (implied barrier)
```



Another OpenMP example



```
1 void mxv_row(int m,int n,double *a,double *b,double *c)
2 {
3     int i, j;
4     double sum;
5
6     #pragma omp parallel for default(none) \
7         private(i,j,sum) shared(m,n,a,b,c)
8     for (i=0; i<m; i++)
9     {
10         sum = 0.0;
11         for (j=0; j<n; j++)
12             sum += b[i*n+j]*c[j];
13         a[i] = sum;
14     } /*-- End of parallel for --*/
15 }
```

```
% cc -c -fast -xrestrict -xopenmp -xloopinfo mxv_row.c
"mxv_row.c", line 8: PARALLELIZED, user pragma used
"mxv_row.c", line 11: not parallelized
```

#pragma omp for ordered *[clauses...]*
(loop region)

#pragma omp ordered

**/*This code is executed in the same order
than the sequential execution (1 thread at a
time)*/**

structured_block
(endo of loop region)

Load balancing



- ❑ *Load balancing is an important aspect of performance*
- ❑ *For regular operations (e.g. a vector addition), load balancing is not an issue*
- ❑ *For less regular workloads, care needs to be taken in distributing the work over the threads*
- ❑ *Examples of irregular workloads:*
 - *Transposing a matrix*
 - *Multiplication of triangular matrices*
 - *Parallel searches in a linked list*
- ❑ *For these irregular situations, the **schedule** clause supports various iteration scheduling algorithms*

Do Scheduling

› SCHEDULE

- Controls how the iterations are assigned to threads.

!\$OMP SCHEDULE (type,[number])

- type: Determines how
- number: Determines how many

The schedule clause/1



schedule (static | dynamic | guided [, chunk])
schedule (runtime)

static [, chunk]

- ✓ *Distribute iterations in blocks of size "chunk" over the threads in a round-robin fashion*
- ✓ *In absence of "chunk", each thread executes approx. N/P chunks for a loop of length N and P threads*

Example: Loop of length 16, 4 threads:

TID	0	1	2	3
no chunk	1-4	5-8	9-12	13-16
chunk = 2	1-2 9-10	3-4 11-12	5-6 13-14	7-8 15-16

The schedule clause/2



dynamic [, chunk]

- ✓ *Fixed portions of work; size is controlled by the value of chunk*
- ✓ *When a thread finishes, it starts on the next portion of work*

guided [, chunk]

- ✓ *Same dynamic behaviour as "dynamic", but size of the portion of work decreases exponentially*

runtime

- ✓ *Iteration scheduling scheme is set at runtime through environment variable **OMP_SCHEDULE***

DO / END DO

Program example

```
dimension A(10),B(10)
```

```
Integer i
```

```
!$OMP PARALLEL
```

```
!$OMP DO
```

```
!$OMP SCHEDULE (STATIC)
```

```
Do i=1,8
```

```
B(i) = A(i) / 2.0
```

```
End Do
```

```
!$OMP END DO
```

```
!$OMP END PARALLEL
```

```
Return
```

```
End
```

1	5
2	6
3	7
4	8

DO / END DO

Program example

```
dimension A(10),B(10)
```

```
Integer i
```

```
!$OMP PARALLEL
```

```
!$OMP DO
```

```
!$OMP SCHEDULE (DYNAMIC,1)
```

```
Do i=1,8
```

```
    B(i) = A(i) / 2.0
```

```
End Do
```

```
!$OMP END DO
```

```
!$OMP END PARALLEL
```

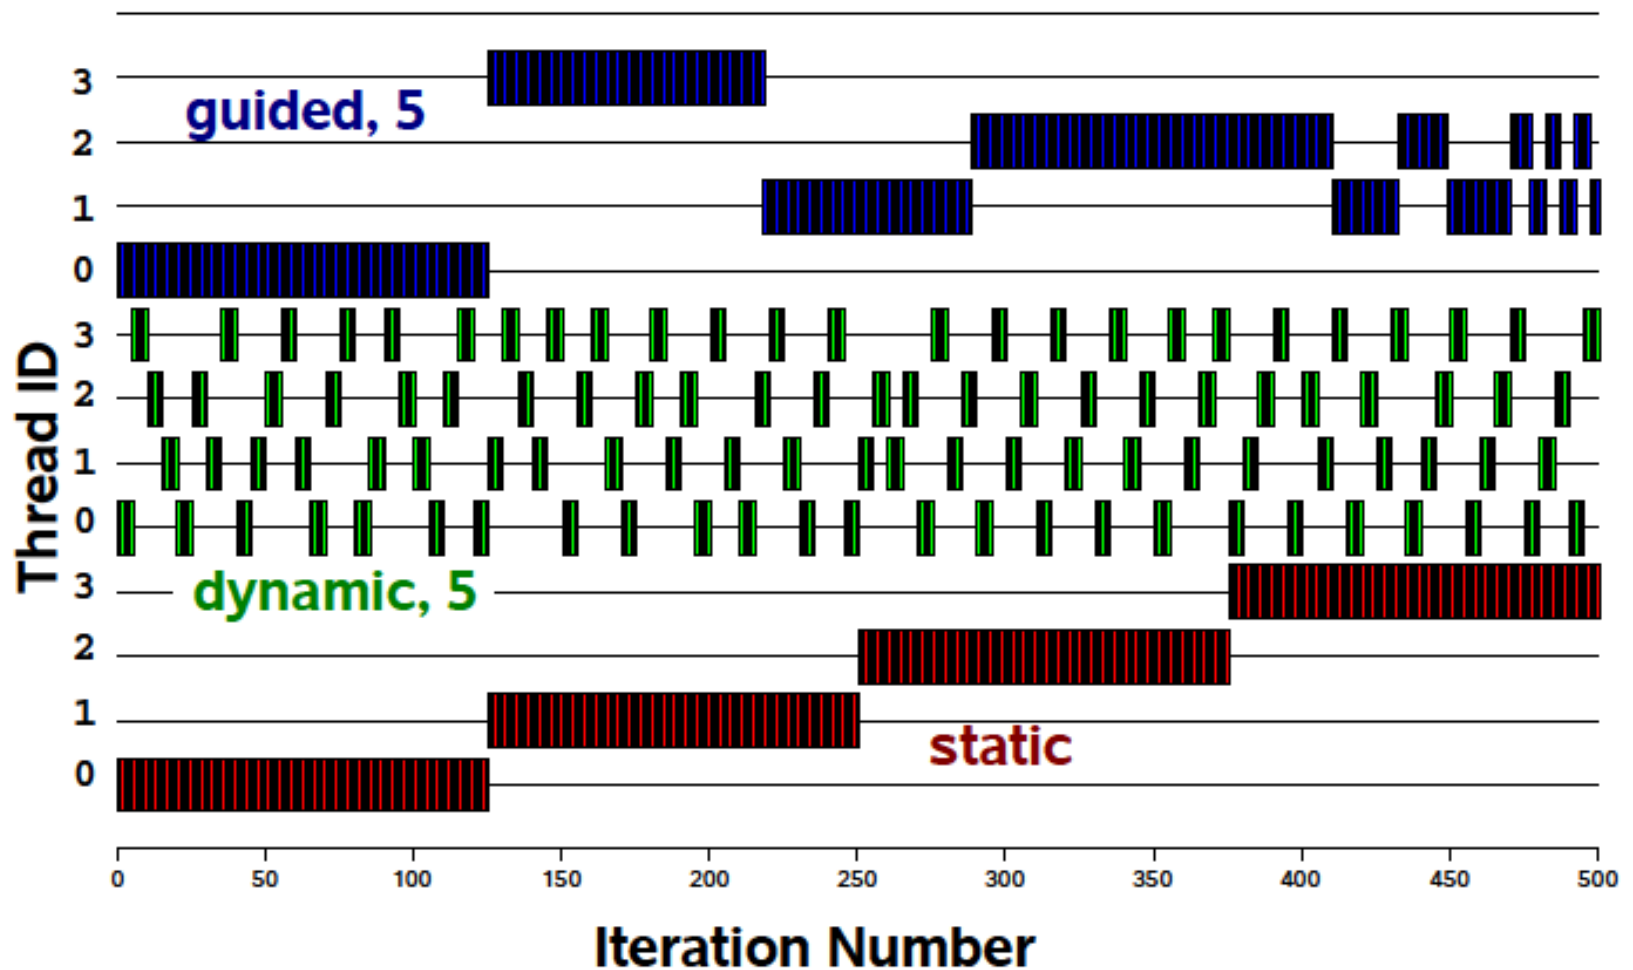
```
Return
```

```
End
```

1	2
3	5
4	7
6	8

The experiment

500 iterations on 4 threads



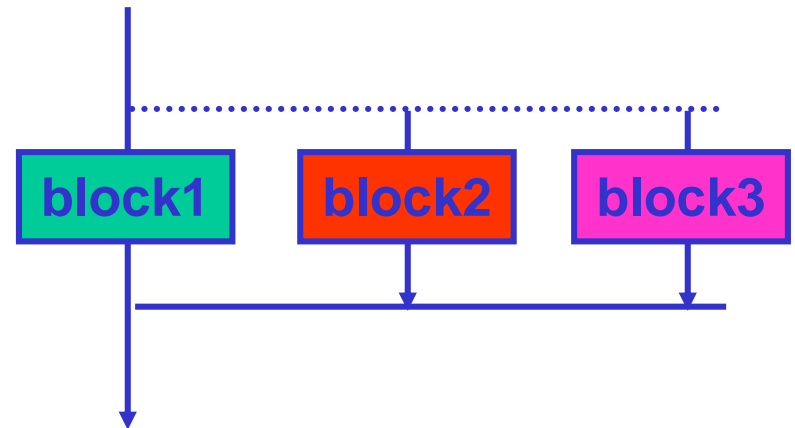
Directives

› SECTIONS / END SECTIONS

- Must be in a parallel region
- Sections are distributed among threads
- Each thread executes a different section
- Allows task level parallelism
- SECTION pragma defines each section

SECTIONS / END SECTIONS

```
!$OMP PARALLEL  
  !$OMP SECTIONS  
    BLOCK1  
    !$OMP SECTION  
      BLOCK2  
    !$OMP SECTION  
      BLOCK3  
  !$OMP END SECTIONS  
!$OMP END PARALLEL
```



Directives

› SINGLE / END SINGLE

- The code included in the single section will be executed by only one thread

!\$OMP PARALLEL

BLOCK1

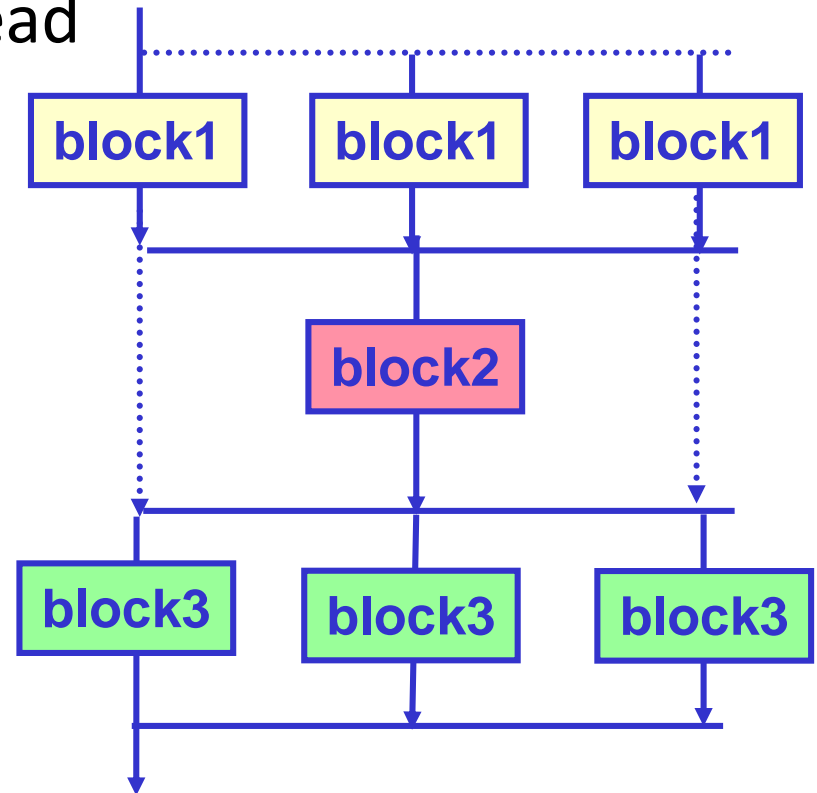
!\$OMP SINGLE

BLOCK2

!\$OMP END SINGLE

BLOCK3

!\$OMP END PARALLEL

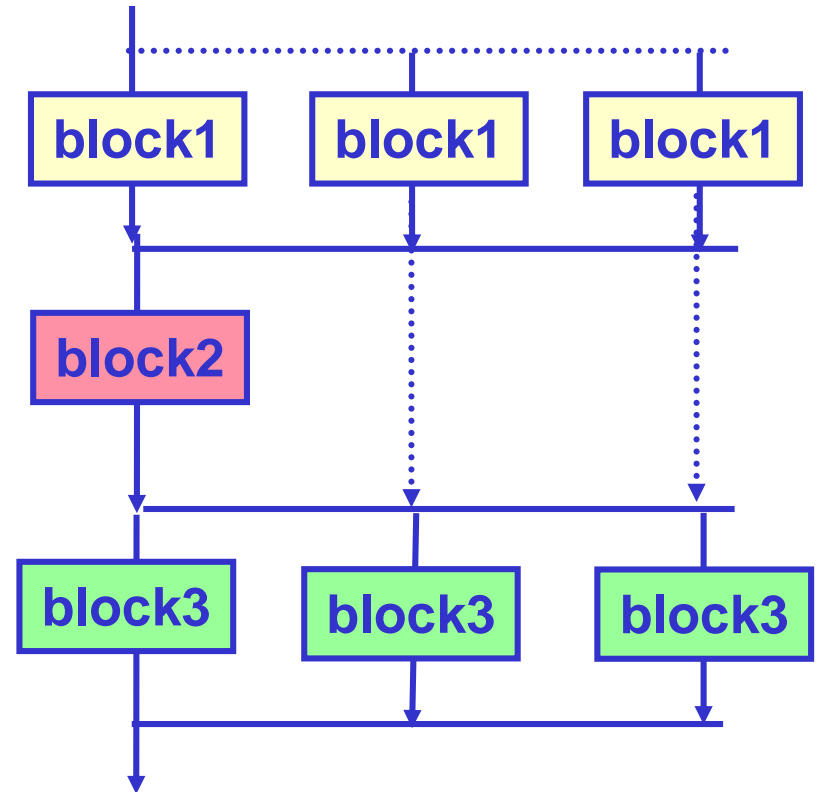


Directives

› MASTER / END MASTER

- Code included in the master section will be only executed by the master thread

```
!$OMP PARALLEL
    BLOCK1
    !$OMP MASTER
    BLOCK2
    !$OMP END MASTER
    BLOCK3
!$OMP END PARALLEL
```



The sections directive - example



```
#pragma omp parallel default(none) \
    shared(n,a,b,c,d) private(i)
{
    #pragma omp sections nowait
    {
        #pragma omp section
        {
            for (i=0; i<n-1; i++)
                b[i] = (a[i] + a[i+1])/2;
        }

        #pragma omp section
        {
            for (i=0; i<n; i++)
                d[i] = 1.0/c[i];
        }
    } /*-- End of sections --*/
} /*-- End of parallel region --*/
```



Short-cuts



```
#pragma omp parallel
#pragma omp for
  for (...)
```



```
#pragma omp parallel for
  for (...)
```

Single PARALLEL loop

```
!$omp parallel
!$omp do
  ...
!$omp end do
!$omp end parallel
```



```
!$omp parallel do
  ...
!$omp end parallel do
```

```
#pragma omp parallel
#pragma omp sections
{ ... }
```



```
#pragma omp parallel sections
{ ... }
```

Single PARALLEL sections

```
!$omp parallel
!$omp sections
  ...
!$omp end sections
!$omp end parallel
```



```
!$omp parallel sections
  ...
!$omp end parallel sections
```

```
!$omp parallel
!$omp workshare
  ...
!$omp end workshare
!$omp end parallel
```



```
!$omp parallel workshare
  ...
!$omp end parallel workshare
```

Single WORKSHARE loop

Barrier/1



Suppose we run each of these two loops in parallel over i:

```
for (i=0; i < N; i++)  
    a[i] = b[i] + c[i];
```

```
for (i=0; i < N; i++)  
    d[i] = a[i] + b[i];
```

This may give us a wrong answer (one day)

Why ?

Barrier/2



We need to have updated all of a[] first, before using a[]

```
for (i=0; i < N; i++)  
    a[i] = b[i] + c[i];
```

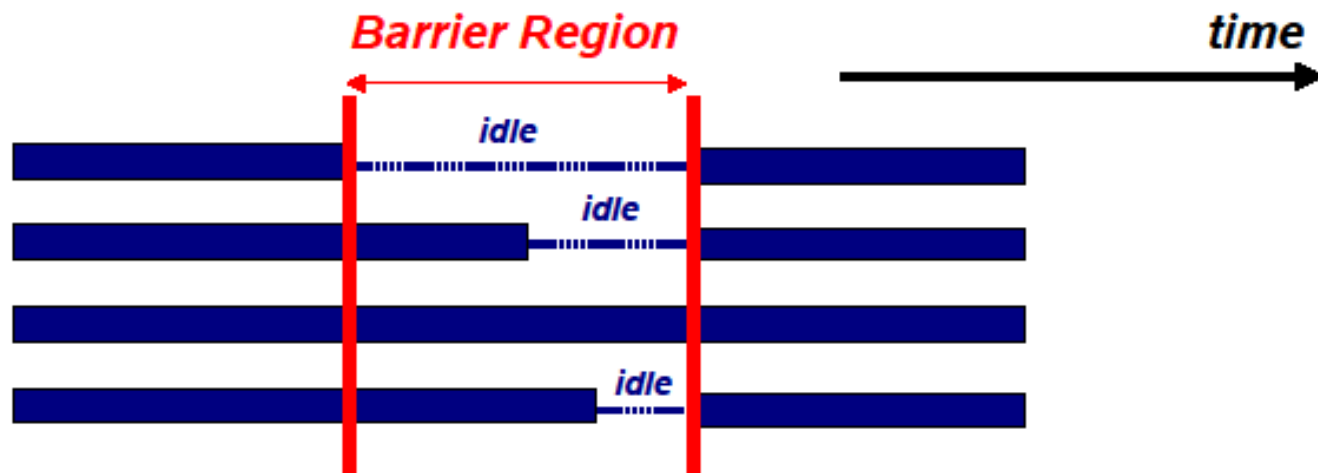
wait !

barrier

```
for (i=0; i < N; i++)  
    d[i] = a[i] + b[i];
```

All threads wait at the barrier point and only continue when all threads have reached the barrier point

Barrier/3



Each thread waits until all others have reached this point:

```
#pragma omp barrier
```

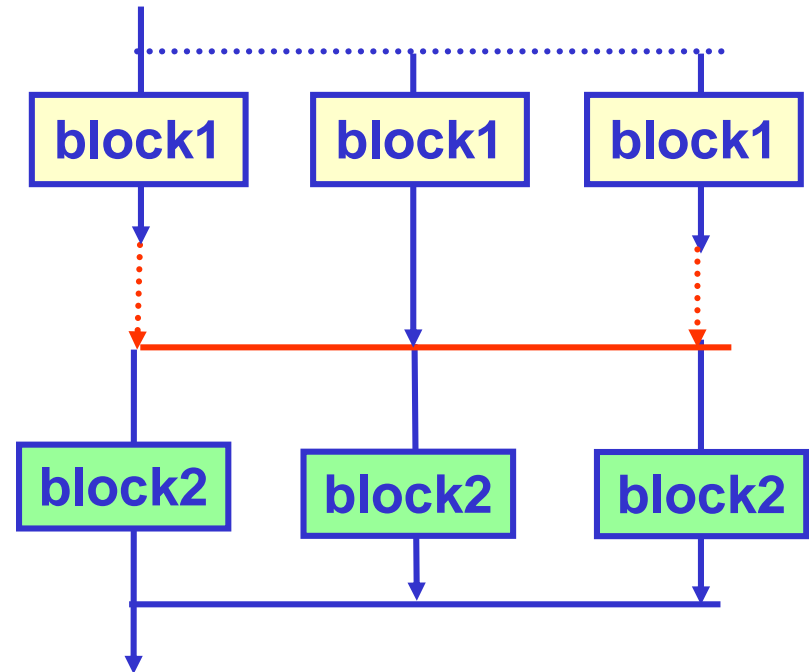
```
!$omp barrier
```

Barrier

› BARRIER

- All threads wait until the last arrives to the barrier

```
!$OMP PARALLEL  
    BLOCK1  
    !$OMP BARRIER  
    BLOCK2  
!$OMP END PARALLEL
```



Critical region/1



If sum is a shared variable, this loop can not be run in parallel

```
for (i=0; i < N; i++) {  
    .....  
    sum += a[i];  
    .....  
}
```

We can use a critical region for this:

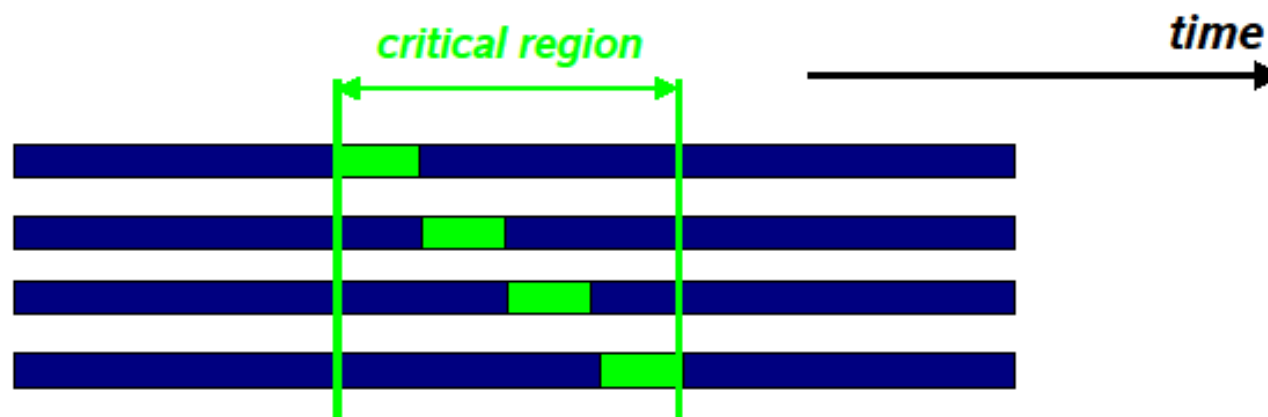
```
for (i=0; i < N; i++) {  
    .....  
    sum += a[i];  
    .....  
}
```

one at a time can proceed

next in line, please

Critical region/2

- ❑ *Useful to avoid a race condition, or to perform I/O (but which still will have random order)*
- ❑ *Be aware that your parallel computation may be serialized and so this could introduce a scalability bottleneck (Amdahl's law)*



CRITICAL / END CRITICAL

```
!$OMP PARALLEL
    BLOCK1
    !$OMP CRITICAL
        BLOCK2
    !$OMP END CRITICAL
    BLOCK3
    !$OMP CRITICAL
        BLOCK4
    !$OMP END CRITICAL
    BLOCK5
!$OMP END PARALLEL
```

Synchronization

- › **There are implicit barriers in:**
 - PARALLEL / END PARALLEL
 - DO / END DO
 - SECTIONS / END SECTIONS
 - SINGLE / END SINGLE
- › **The NOWAIT pragma avoids this barrier**

OpenMP 3.0

› TASK Construct

- The TASK construct defines an explicit task, which may be executed by the encountering thread, or deferred for execution by any other thread in the team.
- The data environment of the task is determined by the data sharing attribute clauses.

OpenMP 3.0

#pragma omp task [*clause ...*]

if (*scalar expression*)

final (*scalar expression*)

untied

default (shared | none)

mergeable

private (*list*)

firstprivate (*list*)

shared (*list*)

structured_block

OpenMP 3.0

- › The TASKWAIT construct specifies a wait on the completion of child tasks generated since the beginning of the current task.

#pragma omp taskwait

```
struct node {  
    struct node *left;  
    struct node *right;  
};  
extern void process(struct node *);  
void traverse( struct node *p ) {  
    if (p->left)  
        #pragma omp task // p is firstprivate by default  
        traverse(p->left);  
    if (p->right)  
        #pragma omp task // p is firstprivate by default  
        traverse(p->right);  
    #pragma omp taskwait  
    process(p);  
}
```



```
#include <stdio.h>
int main()
{
    int x = 1;
    #pragma omp parallel
    #pragma omp single
    {
        #pragma omp task shared(x) depend(out: x)
        x = 2;
        #pragma omp task shared(x) depend(in: x)
        printf("x = %d\n", x);
    }
    return 0;
}
```



OpenMP 3.0

› TASKGROUP Construct

- The TASKGROUP construct specifies a wait on completion of child tasks of the current task and their descendent tasks
- A TASKGROUP region binds to the current task region. The binding thread set of the taskgroup region is the current team

› TASKYIELD Construct

- The TASKYIELD construct specifies that the current task can be suspended in favor of execution of a different task

OpenMP 4.0

- › **Support thread affinity policies** (`proc_bind`, `get_proc_bin`, `OMP_PLACES`)
- › **Support execution on devices** (accelerators)
(`omp_set_default_device`, `omp_get_default_device`, `omp_get_num_devices`, `omp_get_num_teams`, `omp_get_team_num`, and `omp_is_initial_device`)
- › **Reduction clause extended to support user defined reductions**
- › **The concept of cancellation is added**

Run-time Library

› **OMP_SET_NUM_THREADS (SCALAR)**

- Sets the number of threads that will be used in the next parallel region
- Only works if called from a sequential portion of the program

Run-time Library

› **OMP_GET_NUM_THREADS ()**

- Returns the number of threads in the parallel region where it's called
- The default number of threads depends on the application

Run-time Library

› **OMP_GET_THREAD_NUM ()**

- Returns the thread id of the thread that calls it
- Master thread has id 0

Run-time Library

- › omp_in_parallel
- › omp_set_dynamic
- › omp_get_dynamic
- › omp_get_cancellation
- › omp_set_nested
- › omp_get_nested
- › omp_set_schedule
- › omp_get_schedule
- › omp_get_thread_limit
- › omp_set_max_active_levels
- › omp_get_max_active_levels
- › **And so on and so forth**

References

- › Ruud van der Pas “An Introduction Into OpenMP”. Sun Microsystems. iWOMP 2005.
- › OpenMP Architecture Review Board. OpenMP Application Program Interface. Version 2.5. May 2005.
- › OpenMP Application Program Interface. Version 4.0. July 2013
- › www.openmp.org
- › computing.llnl.gov/tutorials/openMP/